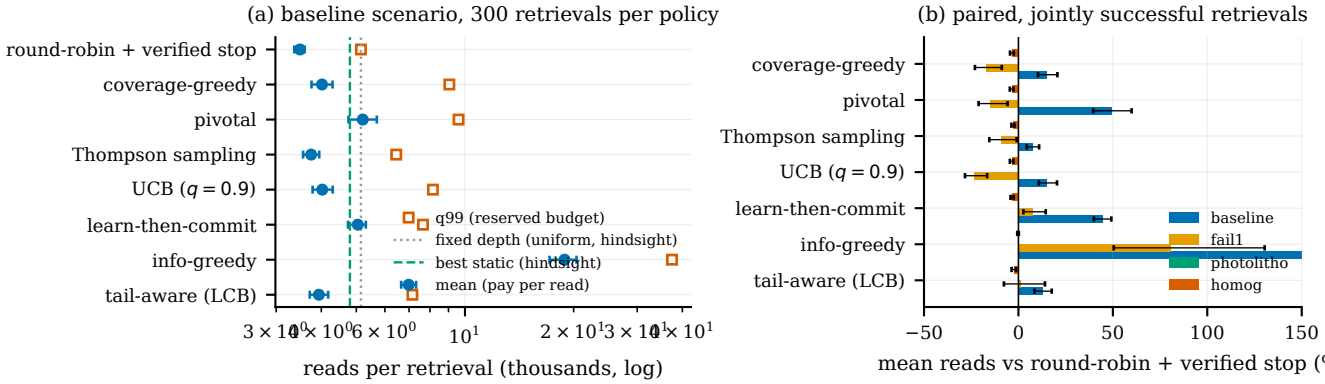
(a) baseline scenario, 300 retrievals per policy
round-robin + verified stop
coverage-greedy
pivotal
Thompson sampling
UCB (q = 0.9)
learn-then-commit
info-greedy
tail-aware (LCB)
q99 (reserved budget)
fixed depth (uniform, hindsight)
best static (hindsight)
mean (pay per read)
reads per retrieval (thousands, log)
(b) paired, jointly successful retrievals
baseline
fail1
photolitho
homog
mean reads vs round-robin + verified stop (%