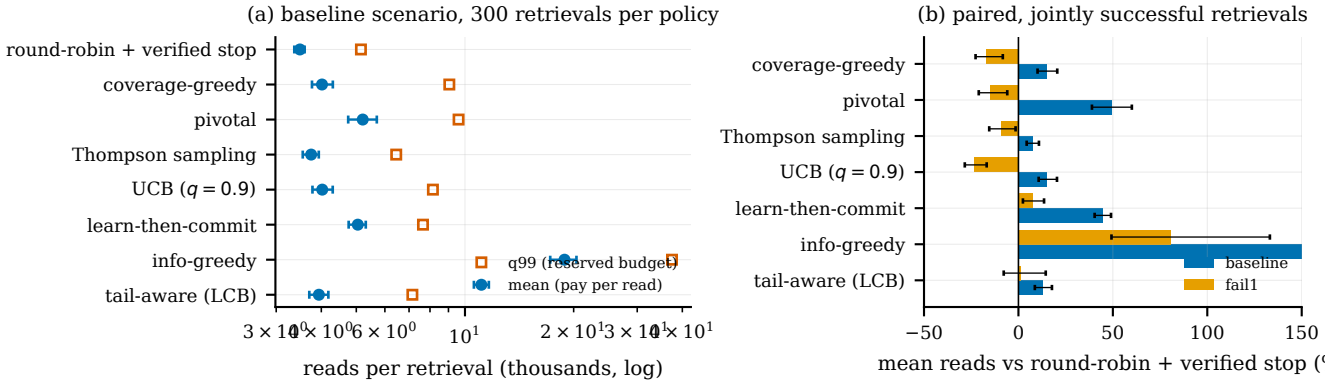
(a) baseline scenario, 300 retrievals per policy
round-robin + verified stop
coverage-greedy
pivotal
Thompson sampling
UCB ($q = 0.9$)
learn-then-commit
info-greedy
tail-aware (LCB)
q99 (reserved budget)
mean (pay per read)
3×10^{0}
6×10^{0}
10^{1}
2×10^{1}
3×10^{1}
4×10^{1}
reads per retrieval (thousands, log)
(b) paired, jointly successful retrievals
coverage-greedy
pivotal
Thompson sampling
UCB ($q = 0.9$)
learn-then-commit
info-greedy
tail-aware (LCB)
baseline
fail1
−50
0
50
100
150
mean reads vs round-robin + verified stop (%